

Bootstrapping Time Series Data

Paul Teetor
Quant Development LLC

CSP 2014
Tampa, Florida


We'll cover a range of bootstrapping procedures today.

- Background on the bootstrap
- Non-parametric: The naïve bootstrap
- Handling dependency: The Moving Block bootstrap
- Honoring a model: Parametric bootstrap
- Balanced approach: The Maximum Entropy bootstrap

When and why do we bootstrap time series data?


- You have some time series data
- But not much data – *whatever “much” means*
- Want to estimate a statistic – *especially a tricky statistic*
- . . . and its confidence interval
- No closed-form solution

Bootstrapping generates *bootstrap replicates* and *replicate statistics*.



Q: How do we get statistic's *conf. interval* from *replicate statistics*?

A: The percentiles of the empirical distribution (histogram) give the confidence interval for the statistic. Cool!



Bootstrapping *time series* data has special challenges.

- Interesting time series are not i.i.d.

We difference the data.

- How do we generate plausible bootstrap replicates?

Several ways. That's what this talk is really about.

- How do we deal with dependency structure?

By choosing the right replication method. Stay tuned.

The bootstrap procedure requires i.i.d. data.

- i.i.d. necessary for resampling with replacement.
- Differencing time series can create i.i.d. data.
- Random walk model, where $\varepsilon_t \sim N(0, \sigma^2)$:

$$y_t = y_{t-1} + \varepsilon_t$$

- Becomes:

$$\varepsilon_t = y_t - y_{t-1}$$

If differences are i.i.d., we can use the *naïve bootstrap*.

Procedure:

- 1) Calculate successive *differences*.
- 2) Repeatedly,
 - 1) Resample the differences with replacement.
 - 2) Sum those differences to construct one replicate time series.
 - 3) Using that time series, calculate one replicate statistic.
- 3) From all the replicate statistics, form the estimate and confidence interval:

Mean of replicate statistics → estimate

Percentiles of replicate statistics → confidence interval

Toy Example

Given time series:

```
[1] 10.00 9.67 9.50 8.66 8.33 7.26 7.48 8.03 8.60 8.44
```

Statistic of interest for given data:

```
[1] 2.74
```

Compute differences:

```
[1] -0.33 -0.17 -0.84 -0.33 -1.07 0.22 0.55 0.57 -0.16
```

Resample the differences with replacement:

```
[1] 0.55 -0.16 -0.84 -0.33 0.22 -0.84 0.22 0.22 0.57
```


Construct one bootstrap replicate (by summing):

```
[1] 10.00 10.55 10.39 9.55 9.22 9.44 8.60 8.82 9.04 9.61
```


Compute one replicate statistic:

```
[1] 1.95
```


Naïve bootstrap example: Stock price, differences, net change




Check: Is it reasonable to assume differences are i.i.d.?



Create replicates by summing resampled differences



Naïve bootstrap example: Statistic of interest is *net change*



Simple implementation in R

```
diffs = diff(price)

HOR = 21

reps = replicate(999,
                 sample(diffs, HOR, replace=TRUE),
                 simplify=TRUE)

reps = apply(reps, 2, cumsum)

outcomes = reps[HOR,]

print(
  quantile(outcomes, prob=c(0.025, 0.975)) )
```

Mean and quantiles of replicate statistics give estimate and conf. int.

```
> summary(outcomes)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-14.430	-1.225	1.120	1.057	3.445	11.540

```
> quantile(outcomes, prob=c(0.025, 0.975))
```

2.5%	97.5%
-6.4120	7.6425


Next problem:
What if the differences are not i.i.d.?

If not, purely random resampling will not capture
the structure of the differences.


Bootstrap replicates will not resemble our data.

Uh oh.

Example: AR(1) time series, *maximum drawdown* statistic



The ACF of this time series reveals
a (simple) dependency.



Moving Block Bootstrap preserves (local) dependency structure.

- Break time series into little blocks.
- Resample the blocks, not individual points –
kind of “random shuffling”, with replacement.
- Within blocks, structure is preserved.
- Works if structure between blocks is (quasi)
i.i.d.

The Moving Block procedure resamples blocks, not points.




Illustration courtesy of <http://www.csp.curtin.edu.au/photos/resample.jpg>

In R, *tsboot* and *boot.ci* together implement a moving block bootstrap.

```
theStatistic = function(x) { . . . }

BLOCK_SIZE = 5          # guess for block size...

mbb = tsboot(ts(xa), theStatistic, R=999,
             l=BLOCK_SIZE, sim="fixed")

replStats = as.vector(mbb$t)

print(summary(replStats))    # for estimate

print(
  boot.ci(mbb, type=c("norm", "basic", "perc")) )
```

Output from *boot.ci*

*** Summary of Replicate Statistics: AR(1) Data, Block Bootstrap

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
6.868	10.350	11.910	11.420	13.390	13.390

*** Confidence Intervals: AR(1) Data, Block Bootstrap

BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS

Based on 999 bootstrap replicates

CALL :


```
boot.ci(boot.out = bout, type = c("norm", "basic", "perc"))
```

Intervals :

Level	Normal	Basic	Percentile
95%	(3.174, 9.245)	(4.240, 9.084)	(8.541, 13.385)

Calculations and Intervals on Original Scale

Sidebar: Normal approx. does not work for maximum drawdown.



What if you have a useful model of your data?

- Example: ARMA, state-space model, or seasonality.
- Model can remove known structure.
- Residuals embody the remaining uncertainty.
- If residuals are i.i.d. time series, we can bootstrap them:

Run the model repeatedly, each time substituting resampled residuals for originals.

For example, let's fit the AR(1) data to a model (with trend term).

*** Fitted AR(1) model:

Call:



```
arima(x = as.ts(xa), order = c(1, 0, 0), xreg = time,  
      include.mean = FALSE)
```

Coefficients:

	ar1	time
-	-0.0329	0.0449
s.e.	0.0995	0.0027

sigma^2 estimated as 2.622: log likelihood = -190.09, aic
= 386.18

Unlike the original AR(1) data, the residuals show no autocorrelation.



Bootstrap residuals by resampling & inserting them into AR(1) process.

If residuals are

$$\varepsilon_1 \dots \varepsilon_T$$


Resample with replacement, giving

$$\varepsilon'_1 \dots \varepsilon'_T$$


And substitute into the AR(1) process:

$$y_t = \delta + \varphi y_{t-1} + \varepsilon'_t$$

Bootstrap replicates will be plausible variations that conform to the model.



Results of bootstrapping AR(1) residuals



If the model's good, it can tighten the final confidence interval.

```
*** Summary of Replicate Statistics: AR(1) Series, Bootstrap of Residuals
Min. 1st Qu. Median      Mean 3rd Qu.      Max.
3.933    7.784    8.767    8.940   10.410   12.330
```

```
*** Confidence Intervals: AR(1) Series , Bootstrap of Residuals
BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
Based on 999 bootstrap replicates
```

CALL :


```
boot.ci(boot.out = bout, type = c("norm", "basic", "perc"))
```

Intervals :


Level	Normal	Basic	Percentile
95%	(5.302, 12.067)	(5.816, 12.507)	(5.118, 11.809)

Calculations and Intervals on Original Scale

Seasonality model suggests resampling across seasons.



Seasonal replicates example: median and 95% conf. bands



Advanced techniques can handle other dependency structures.

<u>Procedure</u>	<u>Structure</u>
Moving block	Stationary; discrete or categorical data
Local bootstrap – <i>Similar to Monte Carlo</i>	Short-range dependence, mild distributional assumption.
Markov bootstrap	Stationary, short-range dependence; discrete or categorical data
Sieve bootstrap	AR(n) models

Is there a middle-ground between naïve bootstrap and full model?

- Naïve is, uh, too naïve.
- Model is often unknown.
- Maximum Entropy bootstrap is alternative.
- Parametric bootstrap of differences.
- Maximum entropy distribution of differences – very mild assumption
- Preserves many properties, including shape, seasonality, even some non-stationarity

Vastly oversimplified outline of maximum entropy bootstrap


- 1) Sort the original data.
- 2) Using sorted data, compute its intermediate points and lower limits for left and right tails.
- 3) Compute the mean of the maximum entropy density within each interval.
- 4) Generate uniform random values on [0,1], and compute sample quantiles at those points.
- 5) Apply to the sample quantiles the correct order to honor the dependence relationships of the observed data.
- 6) Repeat steps 4 and 5 many times (e.g. 999).

Example: This bond market data seems to have structure. Naive bootstrap works poorly.




The maximum entropy bootstrap preserves the gross structure.

Typical Replicates: Market Data, Max. Ent. Boot.



Maximum Entropy Replicates

Obs'ed Sample and Replicates: Market Data, Max. Ent. Boot.



In R, *meboot* package implements the max. entropy bootstrap.

```
library(meboot)

mebOut = meboot(ts(diff(prices)), reps=999)

mebens = mebOut$ensemble

mebens = rbind(prices[1], mebens)

repls = apply(mebens, 2, cumsum)

# 'repls' contains the bootstrap replicates
```

Bootstrapping Time Series Data: Some Limitations

- Problems with sample: non-representative, too small
- Problems from dependency structure: wrong dependency assumption; regime changes; long-term dependency; overlooked completely
- Parametric bootstrap: wrong model; non-stationary (unstable) process, hence unstable parameters
- Problems with certain statistics: “Edge” statistics may require many, many replicates
- Finally, Monte Carlo may be better alternative

Some References

- *An Introduction to the Bootstrap* by Efron and Tibshirani
- *Bootstrap Methods and Their Applications* by Davison and Hinkley
- “The Moving Blocks Bootstrap Versus Parametric Time Series Models”, Vogel and Shallcross, *Water Resources Research* (June 1996)
- “Bootstraps for Time Series”, Bühlmann, *Statistical Science* (2002, No. 1)
- “Maximum Entropy Bootstrap for Time Series”, Vinod and López-de-Lacalle, *J. of Stat. Soft.* (Jan 2009)

Thank you!

Talk materials available at

<http://bit.ly/csp2014-teetor>